

General announcement

- Next Tuesday: Guest lecturer, recording unlikely, and probably highly relevant for projects & exam
 - turn up in person!

Discovery & Security for the Web of Things



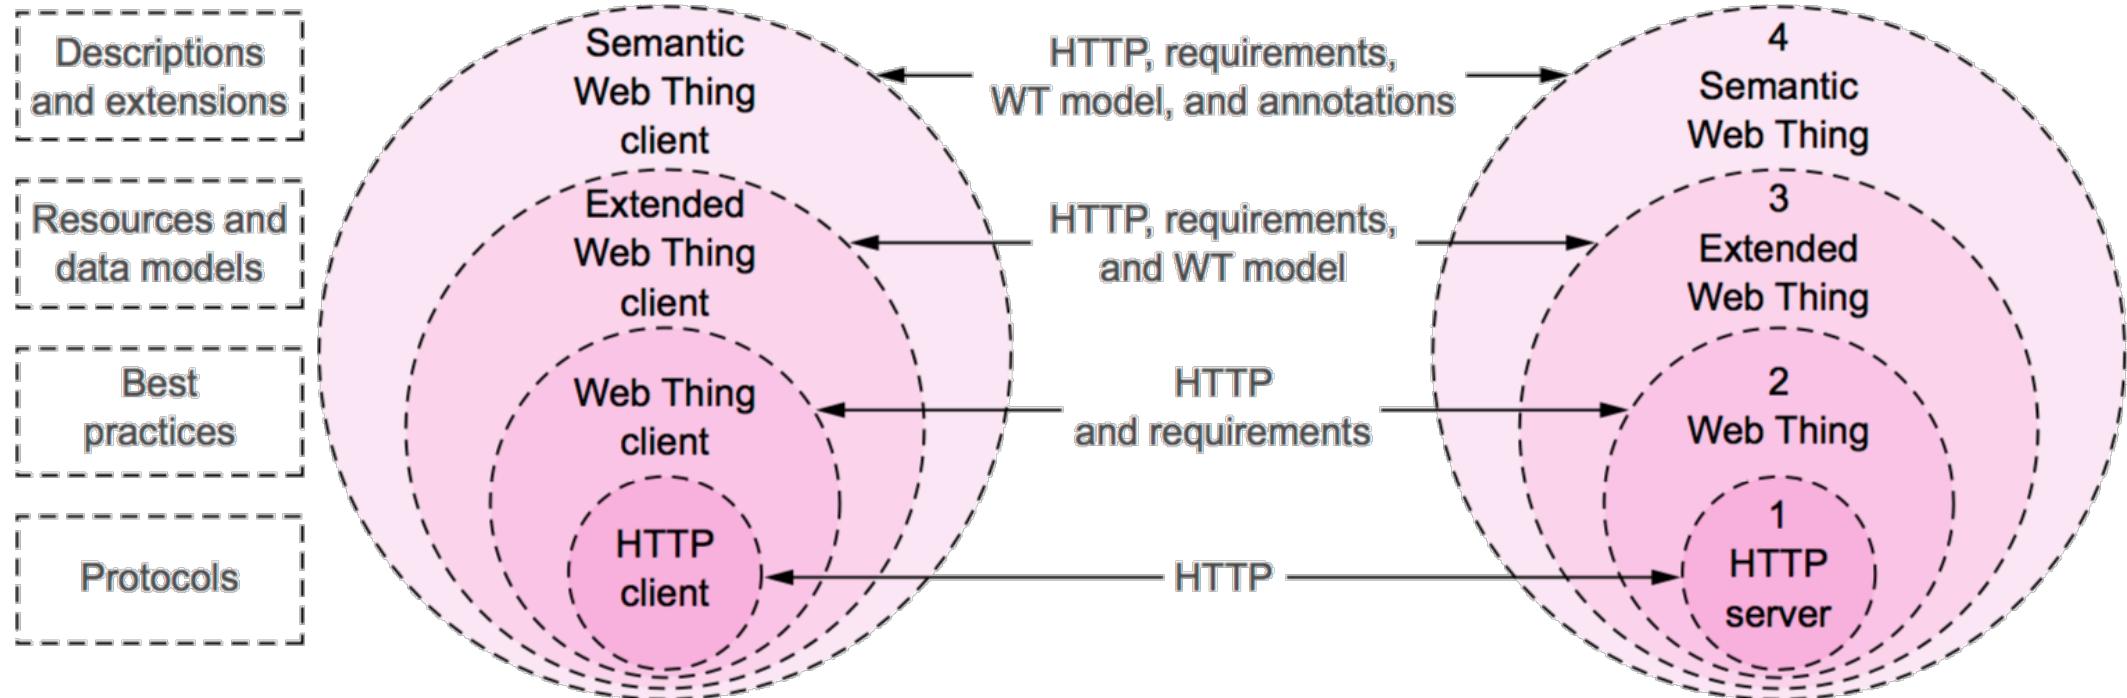
Niels Olof Bouvin

The Challenge of Interoperability

JOKE!

- **What if... there was a *4th* milestone?**
 - “integrate all other groups’Things into your storage system”... due Monday!
- **How would you go about such a task?**
 - collect the URL of each Thing, and
 - inspect each Thing’s API and write custom code for everyone?
 - would *you* want to have to rely on the API documentation that *you* have written?
- **There *has* to be a better way**
 - but we would need to come to a shared understanding and practice for it to work

Levels of specification



- We can vary how deeply we specify our Things
 - it's more work, but it might ease interoperability
- Ideally, it would help make our Things more robust
 - possible to, e.g., verify their API, which is excellent for testing purposes

What do we need?

- A way to discover local Things
- A way to discover the Things' API programmatically
 - sufficiently well to derive required arguments, responses, etc.
- A way to discover Things across the Internet through search engines

Overview

- **Discovering Things**
- **Towards a general model for Things**
- **The Semantic Web of Things**
- **The state of IoT security**
- **Securing and sharing the Web of Things**

Is there any Thing out there?

- **HTTP does not really have *any* discovery mechanism**
 - you are expected to know your destination's address from *somewhere*, often a link
 - once you have reached a page, you can usually navigate to the rest of the site, but this is strictly based on conventions
 - over time, we have grown accustomed to search engines indexing sites, bypassing the need for a proper discovery mechanism, though that frankly is a bit of a hack
- **So, how might we discover which Things are present?**

Discovery services for IP

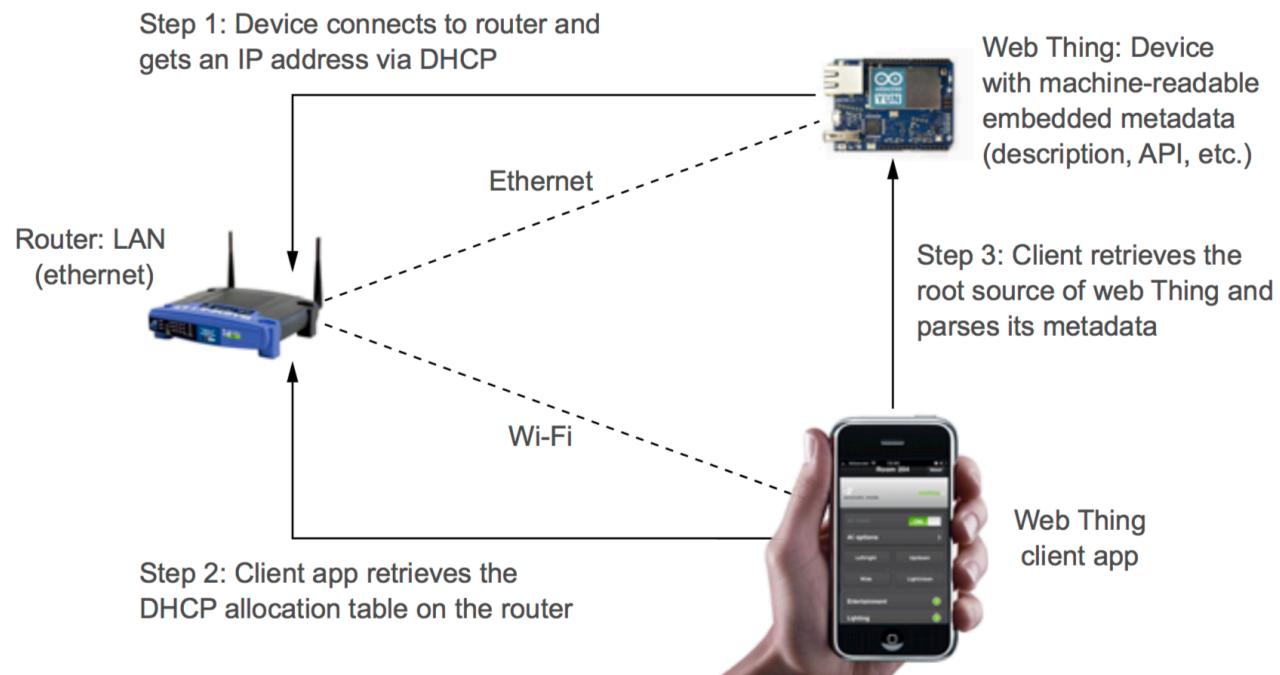
- Discovering new devices on the local network is not a new problem. Often referred to as zero-configuration
 - DHCP
 - UPnP, DLNA
 - mDNS, Bonjour (macOS), Avahi (Linux)
- Usually handled through a broadcast across the local network, where devices can announce or identify themselves
 - to announce/provide services or to search/request for services

Problem...

- These standards are not well supported in Web browsers
 - *used* to just work in macOS Safari, but no more!
- So what do?
 - install an extension?
 - have the sensor register itself somewhere and enquire there?

Possible solutions

- **Have the Things announce themselves**
 - using whatever zeroconf method and discover these announcements with an app
 - stick an QRCode on the device; have it announce its presence as a Bluetooth LE beacon
- **Augment the router to share DHCP addresses in JSON**
 - this is a hack: determining the router's address is going to be guesswork



Overview

- Discovering Things
- **Towards a general model for Things**
- **The Semantic Web of Things**
- **The state of IoT security**
- **Securing and sharing the Web of Things**

Making your Thing indexable

- **The vast majority of all Web sites are discovered through search engines**
 - they crawl sites by following links, index their contents, and make them discoverable through search interfaces
 - not a terribly elegant solution, but it works pretty well
- **If we want our sensor to be discoverable on the Web, it must be indexable by search engines**
 - we could provide HTML pages for all endpoints, making it possible to crawl and index those pages
 - but what about tiny Things? What about JSON and other unHTML formats?
- **Plus, a well-understood structure will ease coding**

IETF RFC5988

→ Request

HTTP 1.1 GET /

Host: MyLittleThingie.io

Accept: application/json

<– Response

200 OK

Link: </model/>; rel="model", </properties/>; rel="properties", </actions/>;
rel="actions", </things/>; rel="things", <<http://model.webofthings.io/>>;
rel="type", </help/>; rel="help", </>; rel="ui"

- A format to communicate relationships between Web resources directly in the HTTP header
- </model/> points to the URL mylittlethingie.io/model, and rel="model" indicates the role of the page
- This is already supported in the <head> tag in HTML, but RFC5988 makes it possible to extend this to all general resources accessible over HTTP

The Web of Things model

- If we are to communicate freely with Things, we need a shared model between them
- This model must encompass all possible Things
 - from a simple product ID in a tag, e.g., on a milk carton
 - over sensors
 - to systems with actuators with complex rules
 - and gateways controlling many other things

The Web Thing Model

Web Thing Clients



Native Mobile App



Web App



Web Thing

Discovers
Web Thing

Create
Actions

Read / Subscribe to
Properties

Control
Non-Web Things

Web Thing

URL: <http://gateway.webofthings.io>

→ **/model** - Name, Description, Tags
- Actions/Properties model

→ **/actions** - ledState
- reboot
- displayText

→ **/properties** Temperature 1,221 | Light 579
Time Online 00:05:59 | Humidity 33.99%

→ **/things** Health Monitor | LilyPad

Non-Web Devices



Bluetooth



ZigBee

Source: Building the Web of Things: book.webofthings.io
Creative Commons Attribution 4.0

- A general model for what a Web Thing actually is

The Web Thing Model: /

- **The root resource of the Thing, which is the virtual representation of the Thing itself**
 - {wt} is the complete URL
- **It can, of course, be altered in the usual fashion**

```
→ REQUEST
GET {wt}
Content-Type: application/json

← RESPONSE
200 OK
Link: <model/>; rel="model"
Link: <properties/>; rel="properties"
Link: <actions/>; rel="actions"
Link: <product/>; rel="product"
Link: <type/>; rel="type"
Link: <help/>; rel="help"
Link: <ui/>; rel="ui"
Link: <_myCustomLinkRelType/>; rel="_myCustomLinkRelType"

{
  "id": "myCar",
  "name": "My super great car",
  "description": "This is such a great car.",
  "createdAd": "2012-08-24T17:29:11.683Z",
  "updatedAd": "2012-08-24T17:29:11.683Z",
  "tags": ["cart", "device", "test"],
  "customFields": {
    "size": "20",
    "color": "blue"
  }
}
```

The Web Thing Model: /properties

HTTP/1.1 200 OK

Content-Type: application/json; charset=utf-8

Link: <<http://model.webofthings.io/#properties-resource>>;

```
{  
  {  
    "id": "temperature",  
    "name": "Temperature Sensor",  
    "values": {  
      "t": 9,  
      "timestamp": "2016-01-31T18:25:04.679Z"  
    }  
    "id": "humidity",  
    "name": "Humidity Sensor",  
    "values": {  
      "h": 70,  
      "timestamp": "2016-01-31T18:25:04.679Z"  
    }  
    "id": "pir",  
    "name": "Passive Infrared",  
    "values": {  
      "presence": false,  
      "timestamp": "2016-01-31T18:25:04.678Z"  
    }  
    "id": "leds",  
    "name": "LEDs",  
    "values": {  
      "1": false,  
      "2": false,  
      "timestamp": "2016-01-31T18:25:04.679Z"  
    }  
  }  
}
```

The Web Thing Model: /actions

- The interface to change the state of various properties of the Thing
- Decouples changing properties directly
 - like accessing an object's state through getter/setters rather than directly modifying a field
 - still no verbs—just nouns

```
HTTP/1.1 200 OK
Content-Type: application/json; charset=utf-8
Link: <http://model.webofthings.io/#actions-resource>;
      rel="type"

[{"id": "ledState",
  "name": "Changes the status of the LEDs"}]
```

```
-> REQUEST
POST {WT}/actions/ledState
Content-Type: application/json
{"ledId": "3", "state": true}

<- RESPONSE
HTTP/1.1 204 NO CONTENT
```

The Web Thing Model: /model

- Collects all information about the Thing
- Based on this, we should be able to know what a Thing, which resources it has, and how they may be accessed and modified, including the type of arguments/results

```
HTTP/1.1 200 OK
Content-Type: application/json; charset=utf-8
Link: <model.webofthings.io>; rel="type"
...
{
  "actions": {
    "link": "/actions",
    "title": "Actions of this Web Thing",
    "resources": {
      "ledState": {
        "name": "Changes the status of the LEDs",
        "values": {
          "ledId": {
            "type": "string",
            "required": true
          },
          "state": {
            "type": "boolean",
            "required": true
          }
        }
      }
    }
  }
},
```

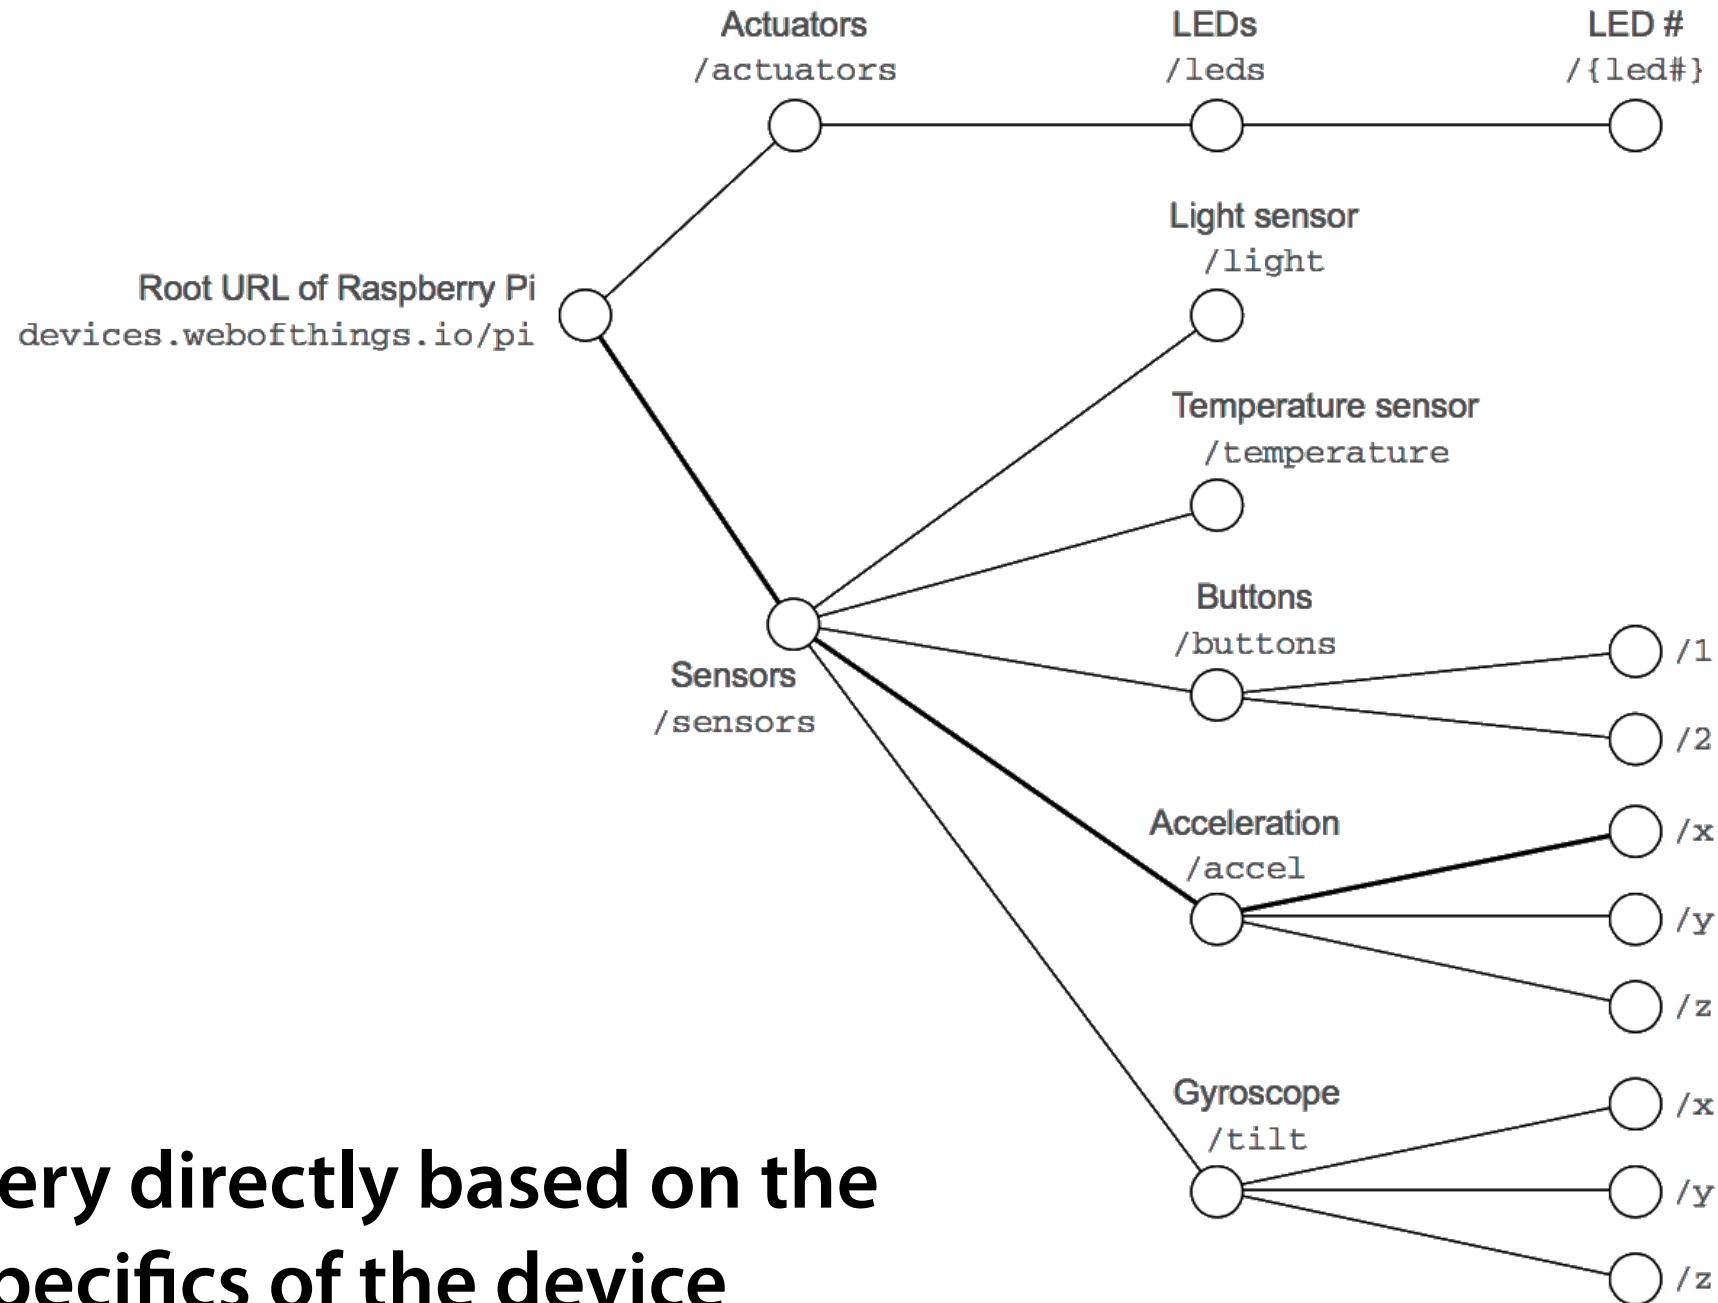
The Web Thing Model: /things

- **The Things that this particular Thing is a gateway for**
 - in this case a webcam and a Hue lightbulb
- **These ‘sub’-Things can be contacted through the Thing**

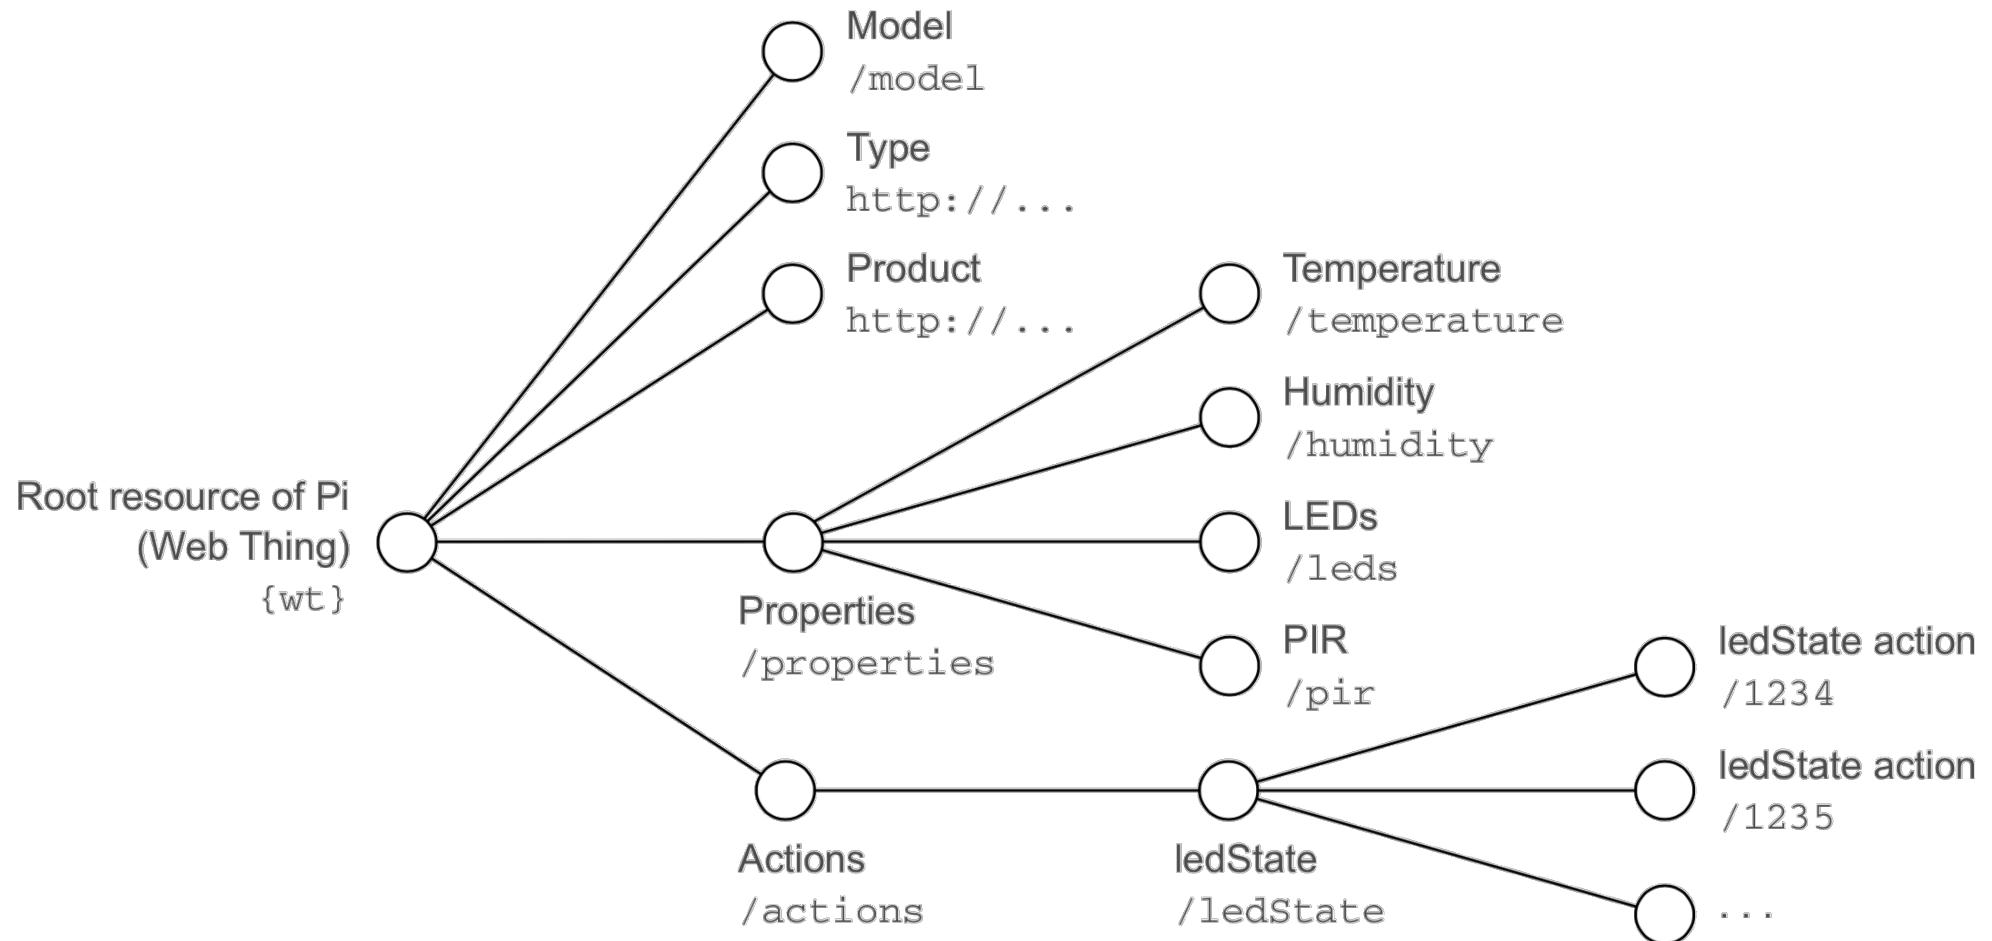
```
HTTP/1.1 200 OK
Content-Type: application/json; charset=utf-8
Link: <model.webofthings.io/things>; rel="meta"

{
  {
    "id": "http://devices.webofthings.io/pi",
    "name": "Raspberry Pi",
    "description": "A WoT-enabled Raspberry Pi"
  },
  {
    "id": "http://devices.webofthings.io/camera",
    "name": "Fooscam Camera",
    "description": "LAN-connected camera."
  },
  {
    "id": "http://devices.webofthings.io/hue",
    "name": "Philips Hue",
    "description": "A WoT-enabled Philips Hue Lamp."
  }
}
```

The old model implementation

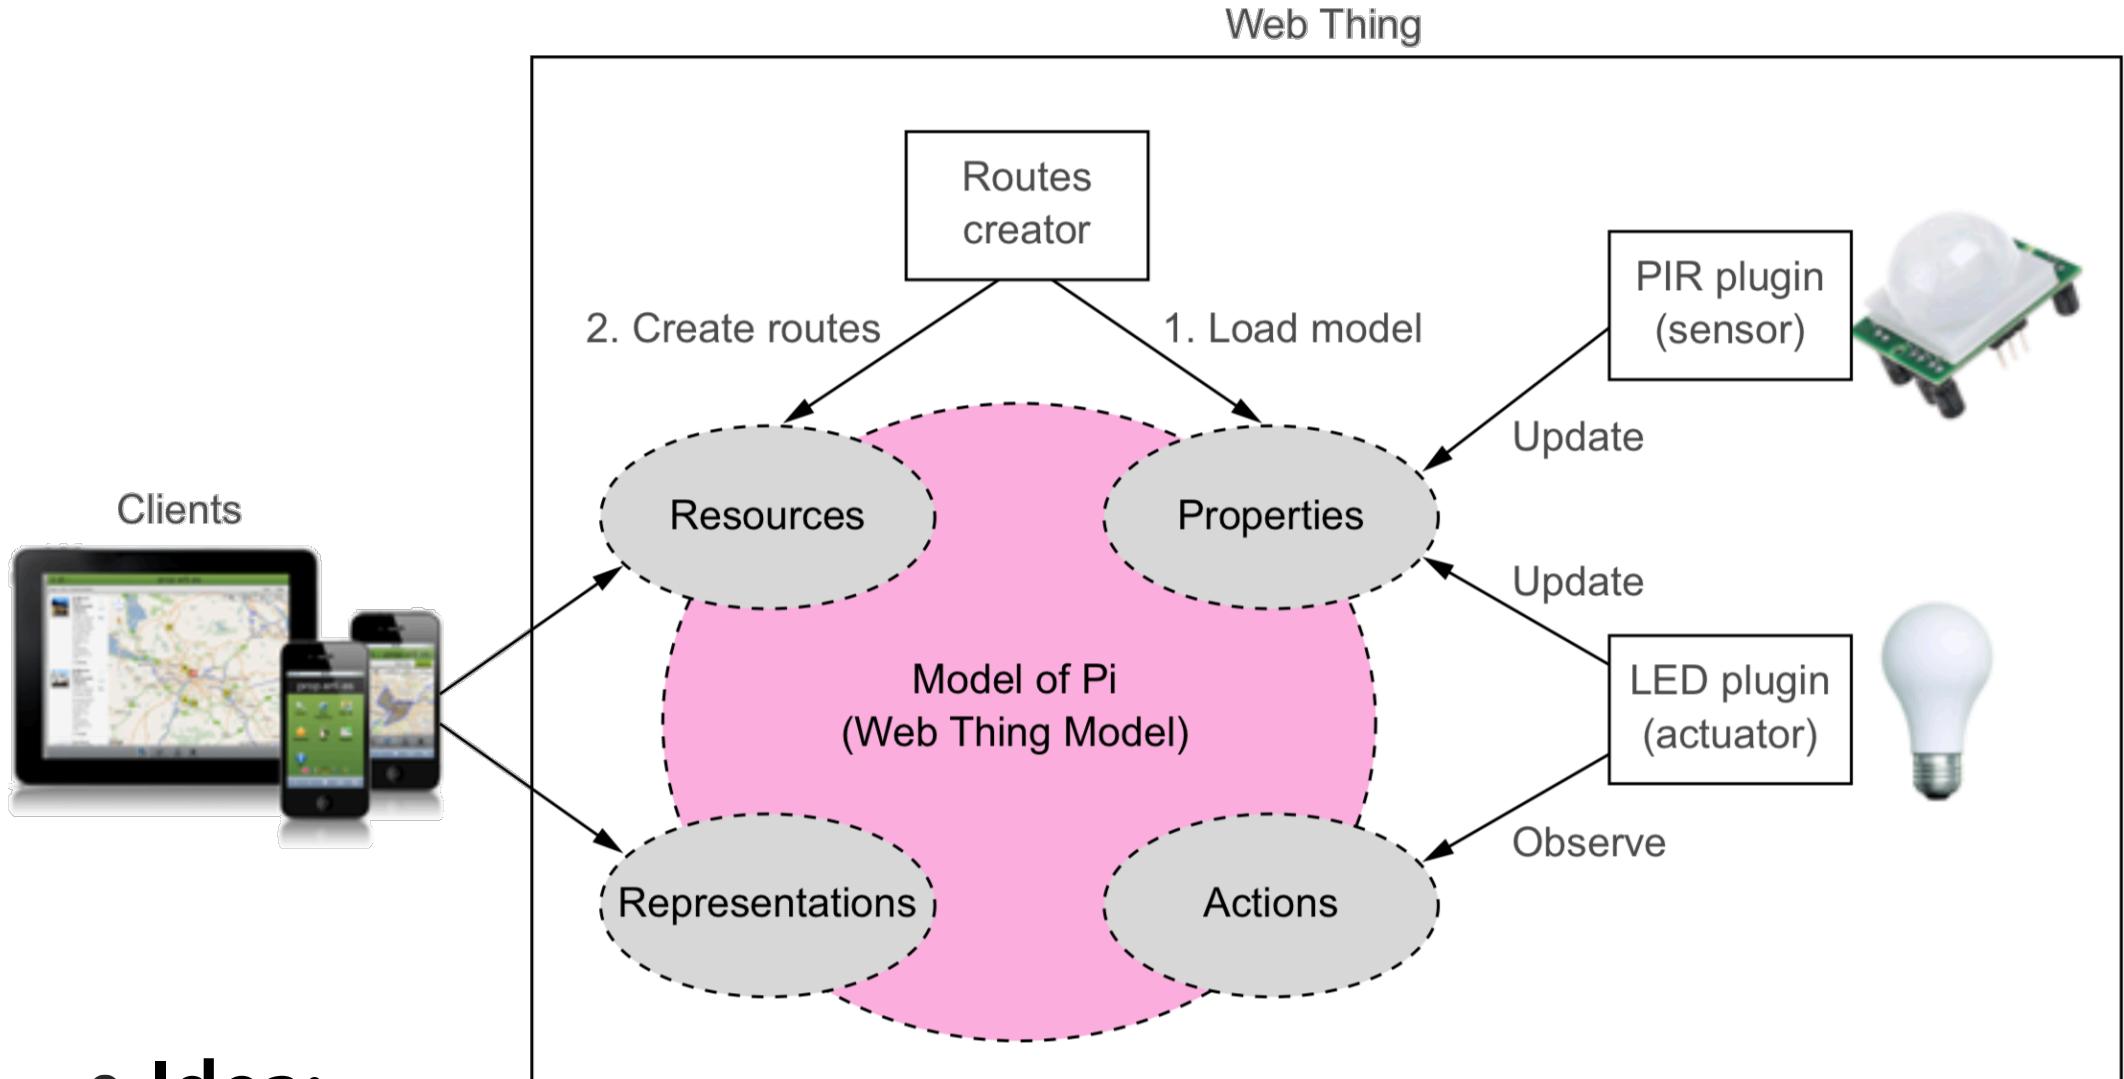


Implementing the WoT model



- A generalised model that can encompass more
 - and, given the right framework, automatically be put to work

The model shall be our guide



- **Idea:**
 - create a comprehensive model of the Thing
 - auto generate routes, etc., based on the model

Auto-connecting code

- It can, of course, not all just be *automagical*
- But, with a sufficient specific model, and prepared code for sensors, it is fairly straightforward
 - see the code in wot-book/chapter8-semantics (available from GitHub)
- The model permits us, by providing a systematic description of the properties and associated actions of the Thing, to create a generic user interface
- Given the identities of the Things, you could, e.g., make a Web page that reported on all temperature sensors in a building

Overview

- Discovering Things
- Towards a general model for Things
- **The Semantic Web of Things**
- **The state of IoT security**
- **Securing and sharing the Web of Things**

The Semantic Web

- An *old* project spearheaded by Sir Tim Berners-Lee, (one of) the inventor(s) of WWW, and W3C's director
- The Web consists of formatted text, as well as other resources
 - while there have been *great* strides in making computers understand text and images, they are still not very good at it compared to humans
- The Semantic Web would have Web resources annotated with meaning, i.e., semantics
 - “this is a person”; “this is an address”; “this is a temperature in Celcius”; etc.
- This would make it far easier for machines to index, and subsequently, make available through search

Challenges for the Semantic Web

- Who authors the metadata?
 - what is their *immediate* benefit?
 - why should we trust their metadata more than their data?
- This has been a recurring problem with the vision of the Semantic Web: in order for it to work as intended, much of the Web must be annotated sufficiently well for machines to extract and understand
 - so far, this has not been going well—for two decades, at least
- Within specific fields of practice and use cases, it can work
 - but it has met with limited success on a greater scale

The reality of the Semantic Web

- If Google et al. do not understand the semantics, there is little point to bother with it
 - The Web Things Model may be a W3C working group's work in progress, but it is not recognised by the search engines
 - *any* business model relying on Google and other search engines to "soon" index specific semantics is overly optimistic
- So, what *do* they index?
 - ordinary Web pages, as you would expect, using heuristics that are top secret to extract meaning (SEO is the attempt to reverse-engineer these secrets)
 - some additional information, added to the Web pages, to make certain elements easier to recognise by the indexer
 - such as identifying things as Products

Providing meaning with annotations

- HTML, as well as JSON can be annotated using various standards and associated *ontologies*

```
<div vocab="http://model.webofthings.io/" typeof="WebThing">
  <h1 property="name">Raspberry Pi</h1>
  <div property="description">
    <p>A simple WoT-connected Raspberry PI for the WoT book.</p>
  </div>
  <p>ID:<span property="id">1</span></p>
  <p>Root URL:<a property="url" href="http://devices.webofthings.io">http://
    devices.webofthings.io</a></p>
  Resources:
  <div property="links" typeof="Product">
    <a property="url" href="https://www.raspberrypi.org/products/raspberry-
      pi-2-model-b/">
      Product this Web Thing is based on.</a>
  </div>
  <div property="links" typeof="Properties">
    <a property="url" href="properties/">
      Properties of this Web Thing.</a>
  </div>
```

If we can't rely on Google...

- **Using semantic markup is not necessarily an exercise in futility, even if Google et al. do not index it properly**
- **It can be used in more specific contexts**
 - writing a general browser side JavaScript library to extract information about Things and provide fancy controls and visualisations
 - targeted towards more specialised search engines that *do* index the semantics

Summary

- **Discovery poses a number of challenges for a Web based architecture for the Internet of Things**
 - it is not well supported from an architectural point of view (wrong level in the stack)
 - the mechanisms in place for Internet scale discovery (search engines) cater to ordinary Web pages rather than Things
 - Semantic annotations may help, but requires adoption by The Powers That Be, and the track record is not exactly promising (which is, admittedly, a Catch-22)
- **Is it, *really*, such a problem?**
 - what are the use cases where you would *want* arbitrary users/systems from across the Internet to connect to your Things?
 - *local* discovery using a zero-configuration system is important, but Internet scale?
 - why not write your own WoT search engine, and make Google point to it instead?

Overview

- Discovering Things
- Towards a general model for Things
- The Semantic Web of Things
- **The state of IoT security**
- **Securing and sharing the Web of Things**

The Internet of Things

- **Myriad interconnected devices, reporting and controlling**
 - many different suppliers
 - many different architectures and systems
 - many different use situations ranging from trivial to absolutely crucial
 - many different actors and agendas
- **What could *possibly* go wrong?**
- **Very early days, yet things are not well**

<http://www.insecam.org/en/bycountry/DK/>

Insecam

Most popular Manufacturers Countries Places Cities Timezones New online cameras FAQ Contacts   

Google Custom Search

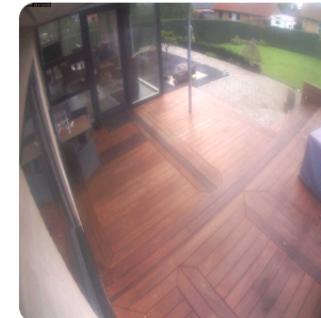


IP cameras: Denmark

« 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 ... 18 »



Watch Foscam camera in Denmark,Viborg



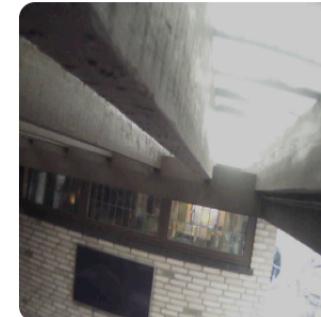
Watch Vivotek camera in Denmark,Tranbjerg



Watch Vivotek camera in Denmark,Silkeborg



Watch PanasonicHD camera in Denmark,Taastrup



Watch Foscam camera in Denmark,Ballerup



Watch Vivotek camera in Denmark,Skanderborg

Types of attacks

- **Denial of Service**
 - removing ability to use a device or a service
 - annoying, if I can't use my IoT toaster; catastrophic, if national power grid is down
- **Surveillance**
 - by the state, by commercial interests, by criminals
- **Intrusion**
 - a root kit on a smart device inside an installation could potentially compromise all devices within network reach

Fear the IoT_Reaper

- One example of a Botnet attacking surveillance cameras, home routers, NAS boxes, etc, last year
 - devices from D-Link, TP-Link, Avtech, Netgear, MikroTik, Linksys, Synology, and others
- Estimates put the number of infected machines around 1-2 mill.

How did we get here?

- Many IoT devices have *very* poor security
 - unencrypted traffic
 - firmware not being patched—either by the manufacturer or the owner
 - best security practices not followed
- They are in homes and companies—*inside* the firewall
 - and can thus act as trojan horses or vectors for attacks
 - as well as surveillance and industrial espionage
 - (this is why *all* communication inside and outside your network should be encrypted)
- This is not bad. This is really, *really* bad

Network level security

- **Challenge: Heterogeneity**
 - strong cryptography may be straightforward to implement on ordinary computers, but what about much more constrained devices?
 - public key infrastructure may be difficult to handle in a large IoT setting
 - gateways can handle part of the burden
- **Centralistic solution simplest, but also a single point of failure**
 - competent actors will act responsible, but that still leaves the rest...
 - and what happens if the company goes under, and the domain is bought by hackers?

Privacy

- **A user's data should belong to the user**
 - unless this can be ensured, the IoT can become the perfect surveillance infrastructure
- **Centralistic solutions easier to exploit**
 - how can the user ensure proper treatment of collected data?
- **Distributed solutions keep data closer to the user (and their control)**
 - but leaves more points to attack
- **EU's General Data Protection Regulation (GDPR) helps**
 - but is also quite complicated and comprehensive, yet it is a start

Identity

- IoT objects must have identities that can be found and authenticated by other services
 - identities can be fixed (5794-118), or fluid (lecture hall for P2P/IoT course)
 - identities can be revealed or hidden (behind authenticated third parties)
- Typical much easier to implement in a centralised system—many challenges remain in a distributed system with ad-hoc connections

Trust

- **How can trust be built?**
- **One thing is the negotiation between devices**
 - based on authentication, negotiation, and observation
- **Another is the trust of users in the IoT**
 - transparency
 - control

Fault tolerance

- **Things *will* go wrong**
- **The system must cope**
 - identify errors and failing sensors
 - choose alternative sensors or services
- **Sometimes systems will come under attack**
 - identify compromised systems
 - route around damage

Summary

- Security and privacy are major requirements for a successful IoT
- So far, there have been plenty of examples of early IoT systems susceptible to attacks
- Clearly, this will have to change
- Industry standards and/or government regulations
- Who owns the data?
 - the generator of the data?
 - the provider of the service?

Overview

- Discovering Things
- Towards a general model for Things
- The Semantic Web of Things
- The state of IoT security
- **Securing and sharing the Web of Things**

The elements of a secure Thing

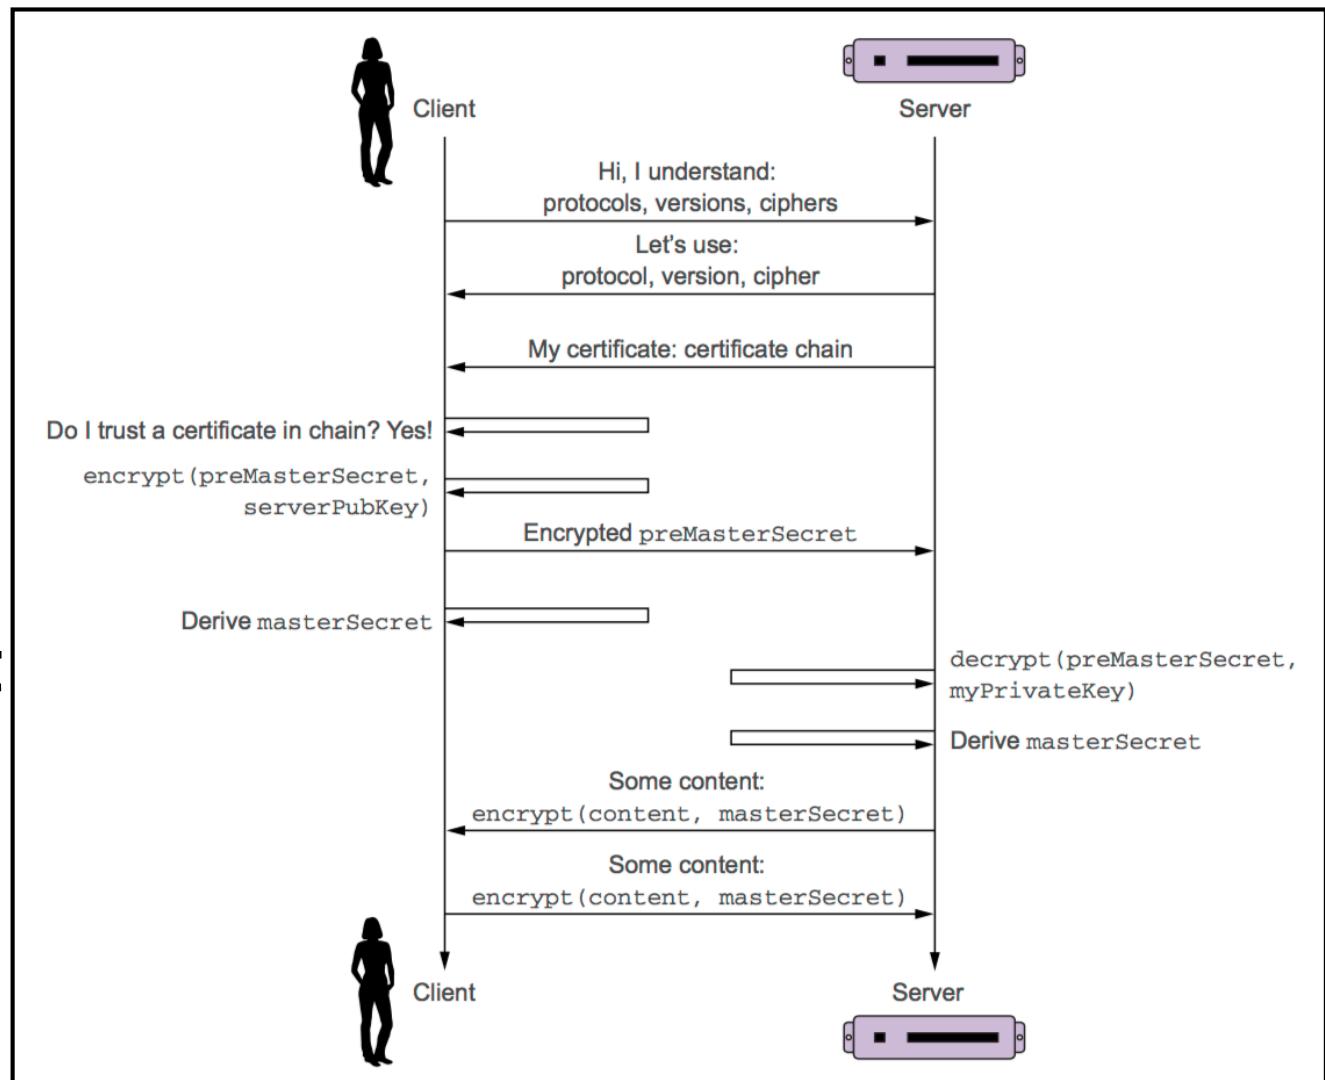
- Encrypted communications
- Authenticated servers
- Authenticated clients
- Secure access control
- Secure software updates

Encrypted communications

- **The basic requirement**
- **Provided that encryption keys are exchanged securely, this should ensure no eavesdropping**
 - assuming best practices are followed, of course
- **Asymmetric and symmetric encryption combined**

Authenticated server

- We need to know the Thing is the *actual* Thing
 - HTTPS and TLS to the rescue
- Keys can be generated locally, and HTTPS support is in Node.js
- <https://letsencrypt.org/> provides free certificates



Authenticated client/user

- **Some IoT devices may not need user authentication**
 - weather stations, public sensors, ...
- **Others certainly require it**
 - cameras, anything with an actuator, anything privacy sensitive

Letting users in

- **Simplest approach: create user profiles with**
 - user names
 - passwords
 - privileges
 - etc
- **Require authentication over secure connection before access is granted**
- **Access, once granted, handled through a token**
 - generated by the server
 - exchanged in headers between client and server

OAuth

- Rather than having users remember yet another password (and having to store that securely), let users connect using preexisting identities
- The user authenticates themselves to a known service, and that service then authorises access to their API for that user from your server
- Not a perfect system—users have been fooled by phishing attacks with sites purporting to be, e.g., Google Docs or nemID requesting authorisation
 - though this is no different than ordinary phishing attacks

Roles in OAuth

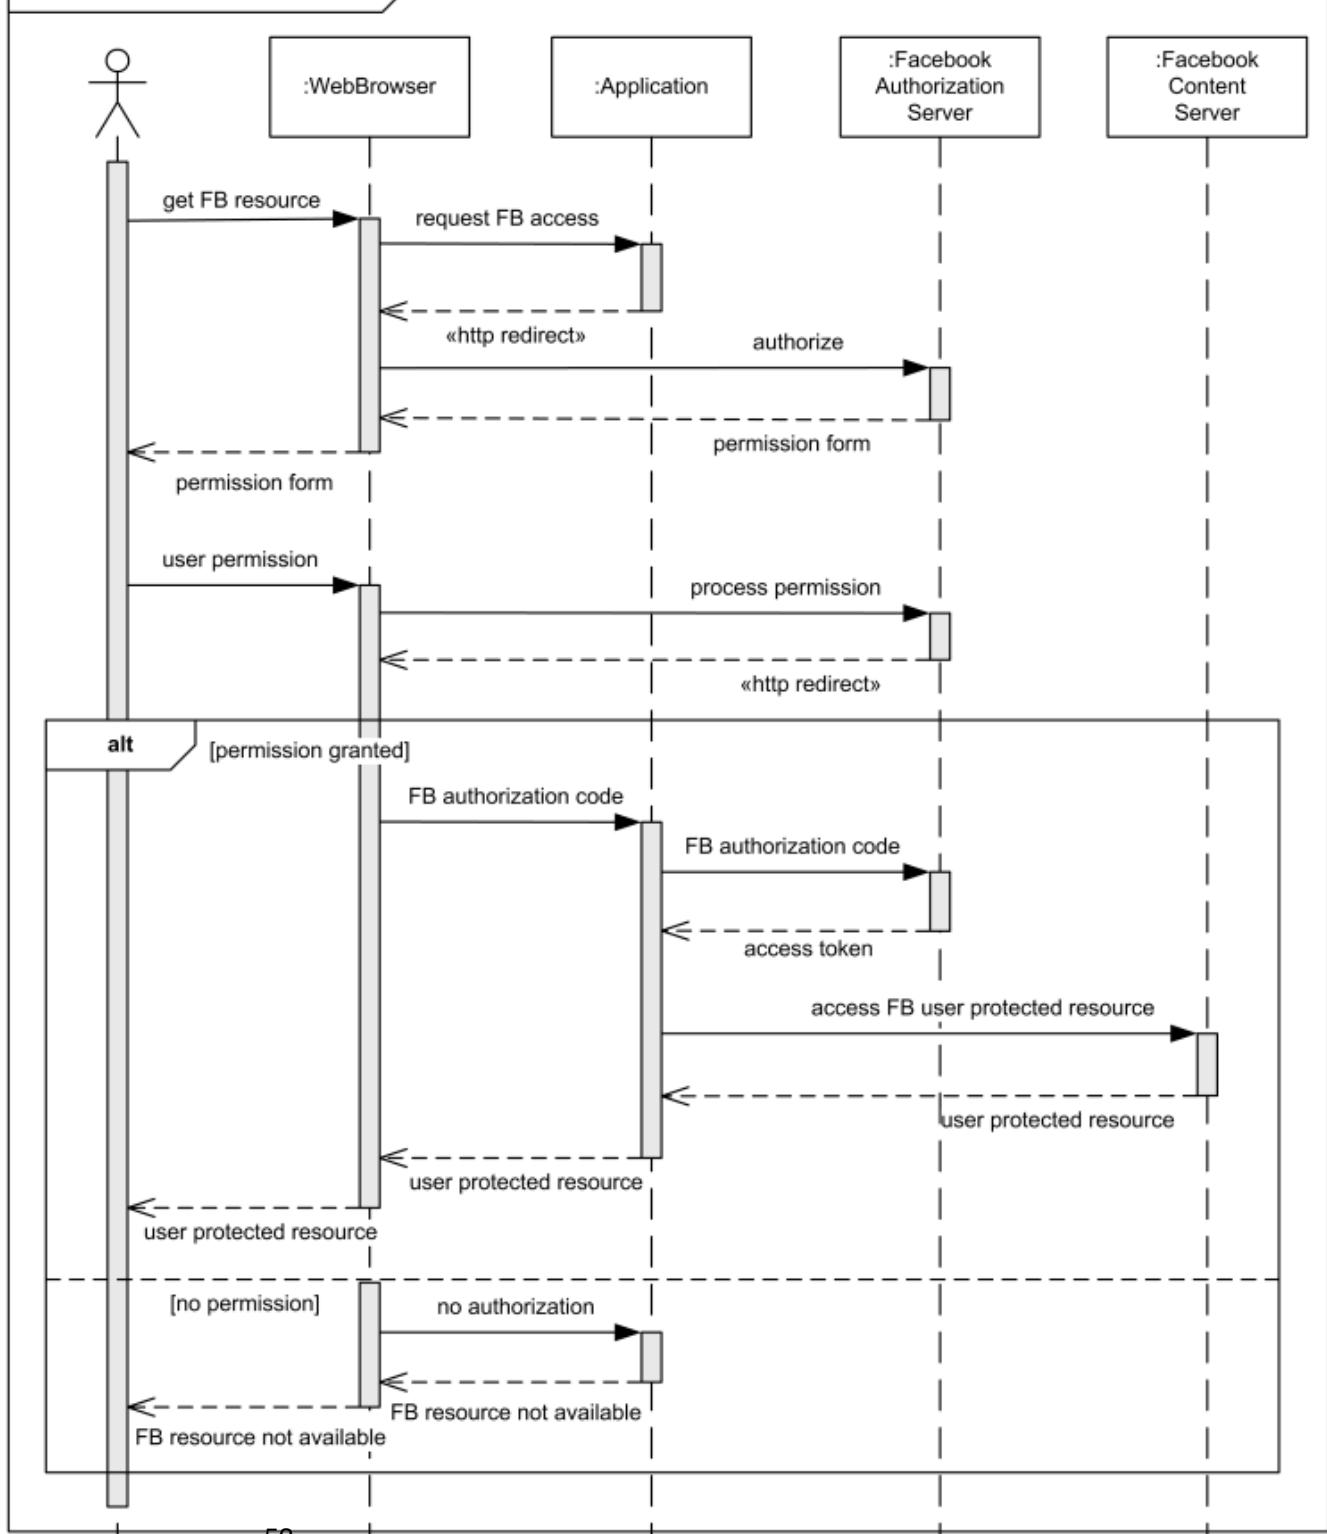
- **The application**
 - the application that needs access to some part of the user's account for, e.g., identity
- **The resource server**
 - provides the API for accessing the user's account
- **The authorisation server**
 - handles the interaction with the user granting/denying access (i.e., login to Twitter)
- **The resource owner**
 - the user, who is granting/denying access to part of their account at the resource server

Requisites

- **The application must be registered with the authorisation server, which provides**
 - client id (this can be public information)
 - client secret (this cannot)
- **The application must provide an redirection URL**
 - which must be secure, e.g., HTTPS

- The interaction between a user, an application, and Facebook authorisation & content servers

- Tokens exchanged through redirects
- Given the access token, the application can access the Facebook API

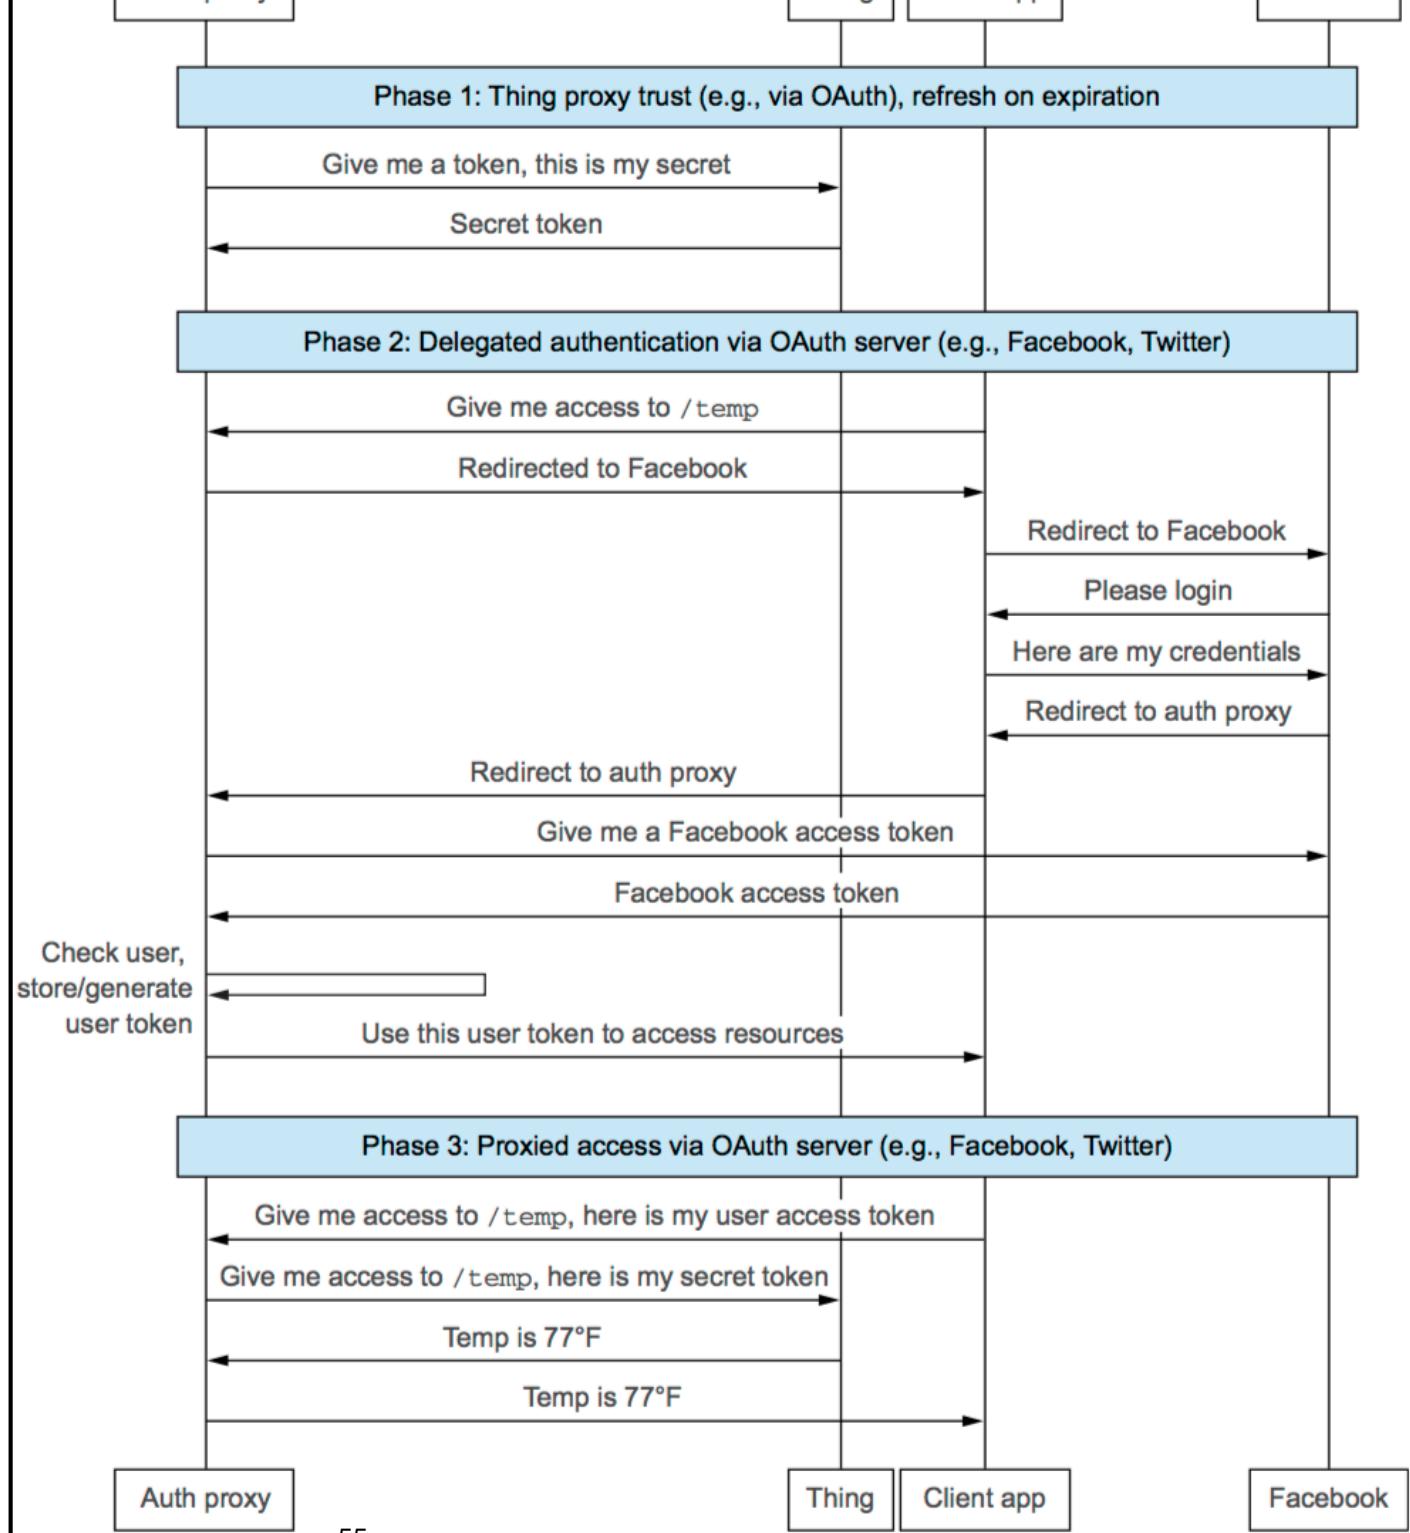


The Social Web of Things

- Using OAuth helps with authorisation and identities of users
- Having each and every Thing maintain lists of users is a bit cumbersome though...
- An Authentication Proxy could handle this interaction
 - registering all the Things
 - handling access to the authorisation servers through OAuth
- One Authentication Proxy for, e.g., a building, a company, or a home

SWoT flow

- Using the token generated, the user can then access the protected resources
 - in *Fahrenheit*?



User identity is not enough

- User roles must also be defined, also known as Access Control Lists
- Some users can be administrators, others cannot or should not

So... what if there is a bug?

- Software contains bugs, and so will IoT devices
- Providing a secure mechanism to push software updates out to devices becomes crucial
 - if the software is not updated, security holes may not be patched, and new features cannot be added
 - if the update mechanism is compromised, the device can be loaded with malware by criminals
 - one solution is an App Store with automatic downloads
- resin.io provides such a service
 - based on git, node.js and/or Docker
 - and it's free, if you have no more than 10 devices

Security is tricky

- **Security is not an end goal—it is an ongoing process**
 - all software contains bugs, including security software, so we must adapt over time
 - the web stack is the most used in the world, so its security will receive much scrutiny
- **Many IoT devices provide terrible security, which puts only the devices and their users at risk, but the entire Internet, if these devices are weaponised into botnets**
- **Therefore: do better, be careful, follow best practices**